

CSE103 Introduction to Algorithms

TD2: Sorting

March 11th, 2021

In this problem sheet you will be asked to implement all the algorithms discussed in Lecture 2 and you will also be introduced to some new material: three more sorting algorithms (*insertion sort* (Exercise 3), *counting sort* (Exercise 6) and *radix sort* (Exercise 7)) as well as the notion of *stability* (Exercise 8). It is suggested that you start with the exercises concerning the new material, so that, if you have any questions about them, you may ask them during the lab (the only exercise on new material that also requires Lecture 2 material is Exercise 8, so you will have to do that last).

You are requested to upload your work on Moodle as a single text file (there is a “TD2” assignment for this). As for the first lab, *this lab is not assessed*, your work will *not* be graded, it is just needed to have an idea of what/how you are doing. Please divide the file you submit in sections, one for each exercise, making the separation evident and explicitly writing the number of the exercise, *in the natural order from 1 to 8* (even if you do the exercises in a different order). For example:

```
#####
# Exercise 1

some Python code

#####
# Exercise 2

some more Python code

...
```

Let us stress once more that CSE103 is about learning and analyzing algorithms, not Python. To make this more conspicuous, from now on you will be *forbidden* to use Python-specific shortcuts in your code. That is, you are requested to use only the subset of Python which is common with any other programming language (essentially).

More specifically, for this problem sheet you will only use:

- function declarations (`def`), `if...elif...else` statements and `while` loops, as well as `for` loops with the `range(m,n,k)` construct restricted to the case in which `m`, `n` and `k` are *non-negative* integer variables or constants (not all three of them must be present, of course);
- basic datatypes like `int`, `bool` and `str`, basic arithmetic on integers (+, -, *, //, %), as well as increment/decrement assignments (`+=`, `-=`), standard comparison operators (`==`, `>`, `>=`, etc.), Boolean operations (`and`, `or`, `not`);
- lists are limited to the basic functionalities:
 - creation of a list of length `n` (either variable or constant): `l = [0] * n`;
 - accessing an element of a list `l[i]` where `i` is a *non-negative* integer variable or constant;
 - concatenation (+) and list length (`len`).

The only Python-specific shortcuts you are allowed to use are:

- one-line swap instructions `a, b = b, a`;
- declarations of default values for function arguments: `def f(a, b=0):`, so that calling `f(2)` is the same as calling `f(2, 0)`.

You will also be allowed to use list comprehension to succinctly initialize a random list.

Nothing else is authorized. In particular, you **may not** use:

- list slicing or any list method (`append`, `insert`, `remove`, `pop...`);
- any other complex type (sets, dictionaries...) or any built-in function (`sum`, `max`, `log...`).

Exercise 1: binary search

The binary search algorithm takes as input an element a and an ascendingly sorted list l , and returns True if a is in l , or False otherwise, as follows:

- if l is empty, return False;
- otherwise, compare the “middle” element $l[m]$ of l with a ;
- if they are equal, return True;
- otherwise:
 - if $l[m] > a$, recursively apply the algorithm to the list $l[0], \dots, l[m-1]$;
 - if $l[m] < a$, recursively apply the algorithm to the list $l[m+1], \dots, l[n-1]$, where n is the length of l .

1. Implement binary search.
2. Modify (and implement) the algorithm so that it returns a position in l where a is found, or -1 if a is not in l . Be careful: if a is found, the position returned must be the one in the original input list l , not the position of a in the sublist analyzed at the current recursive call. For example, searching for 3 in $[0, 2, 3]$ will start by comparing 3 with 2, then recurse on the sublist $[3]$ and find 3 at position 0 in that list, but the position to be returned is 2!

Solution.

1. Here's a possible implementation:

```
def binSearch(a, l):
    if l == []: # not found
        return False
    mid = len(l) // 2
    if l[mid] == a: # found!
        return True
    if l[mid] > a: # construct left half...
        h = [0] * mid
        for i in range(mid):
            h[i] = l[i]
        return binSearch(a, h) # ...and search there
    else: # construct right half...
        h = [0] * (len(l) - mid - 1)
        for i in range(len(h)):
            h[i] = l[i + mid + 1]
        return binSearch(a, h) # ...and search there
```

2. It is enough to keep track of the current “offset”, that is, the position in the original list marking the beginning of the sublist which is the input of the current recursive call:

```
def binSearchPos(a, l, offs=0):
    if l == []:
        return -1
    mid = len(l) // 2
    if l[mid] == a:
        return offs + mid
    if l[mid] > a:
        l1 = [0] * mid
        for i in range(mid):
            l1[i] = l[i]
        return binSearchPos(a, l1, offs)
    else:
        l1 = [0] * (len(l) - mid - 1)
        for i in range(len(l1)):
            l1[i] = l[i + mid + 1]
        return binSearchPos(a, l1, offs + mid + 1)
```

```

else:
    l1 = [0] * (len(l) - mid - 1)
    for i in range(len(l1)):
        l1[i] = l[i + mid + 1]
    return binSearchPos(a, l1, offs + mid + 1)

```

Notice how we used the Python feature allowing us to declare a default value for an input, so that when we write `binSearchPos(a,l)` it is the same thing as `binSearchPos(a,l,0)`. This Python “trick” is quite handy and is allowed. In languages not having this feature, we would have to declare the above 3-parameter function as an auxiliary function like `binSearchPosRec` and then define the actual function as

```

def binSearchPos(a,l):
    return binSearchPosRec(a,l,0)

```

Exercise 2: merge sort

Recall the MERGE problem:

MERGE:

input: two lists l_1, l_2 sorted ascendingly;
output: the list $l_1 + l_2$ (concatenation) sorted ascendingly.

We know that there is an easy efficient algorithm for MERGE: starting from the first elements of l_1 and l_2 , progressively build the result by comparing the current elements of l_1, l_2 and selecting the smallest, until no element is left in one of the two lists, at which point we just append the rest of the other list to the result.

The merge sort algorithm takes as input a list l and returns the ascendingly-sorted version of l as follows:

- if l is empty or a singleton, return l ;
- otherwise, let l'_1, l'_2 be the lists resulting from recursively applying the algorithm to the left and right half of l , respectively;
- return the list obtained by merging l'_1 and l'_2 .

Implement the above algorithm for MERGE and then, using it, implement merge sort. For efficiency reasons, it is better if you do *not* use recursion in the implementation of the MERGE algorithm; just use a loop instead. (You will use recursion for merge sort itself, of course)

Solution. Here is a possible implementation:

```

def merge(l1, l2):
    n1 = len(l1)
    n2 = len(l2)
    l = [0] * (n1 + n2)
    i1 = 0
    i2 = 0
    i = 0
    while i1 < n1 and i2 < n2:
        if l1[i1] <= l2[i2]:
            l[i] = l1[i1]
            i1 += 1
        else:
            l[i] = l2[i2]
            i2 += 1
        i += 1

```

```

    if i1 < n1:
        for i1 in range(i1, n1):
            l[i] = l1[i1]
            i += 1
    elif i2 < n2:
        for i2 in range(i2, n2):
            l[i] = l2[i2]
            i += 1
    return l

def mergesort(l):
    n = len(l)
    if n <= 1:
        return l
    h = n // 2
    l1 = [0] * h
    l2 = [0] * (n-h)
    for i in range(n):
        if i < h:
            l1[i] = l[i]
        else:
            l2[i-h] = l[i]
    return merge(mergesort(l1), mergesort(l2))

```

Exercise 3: insertion sort

Merge sort is a divide and conquer algorithm: it is based on the idea that a list l of length n may be sorted by merging the sorted version of the two halves of l , of length $\lfloor n/2 \rfloor$ and $\lceil n/2 \rceil$. If, instead of splitting l in half, we systematically split it in the singleton containing the first element of l and the rest of the list (of length $n - 1$), we obtain a recursive, non-divide-and-conquer algorithm known as *insertion sort*:

```

def insertionSort(l):
    n = len(l)
    if n <= 1: # if empty or singleton, already sorted
        return l
    else:
        # l1 = l without first element
        l1 = [0] * (n-1)
        for i in range(1, n):
            l1[i-1] = l[i]
        # insert first element into sorted version of l1
        return insert(l[0], insertionSort(l1))

```

where `insert(a, l)` is a function taking an element a and an ascendingly sorted list l , and returning the ascendingly sorted list obtained by inserting a in l at the “right” position. That is, `insert(a, l)` has the same result as `merge([a], l)`, where `merge` implements MERGE as in Exercise 2. For example, `insert(3, [0, 0, 2, 4, 5])` returns `[0, 0, 2, 3, 4, 5]`.

The interest of insertion sort is that, unlike merge sort, it is easily implementable in place and without recursion, just with loops:¹

- let n be the length of the input list l ;

¹Another reason why insertion sort is useful is that it is employable as an *online algorithm*: the input l does not need to be completely known in advance, the list may be sorted while its elements arrive one after the other, asynchronously.

- if $n \leq 1$, exit (there is nothing to do);
- otherwise, for i going from 1 to $n - 1$, insert each element $l[i]$ in the sublist $l[0], \dots, l[i - 1]$. This is done in place: progressively swap the element $a := l[i]$ with the elements of $l[0], \dots, l[i - 1]$, starting from the right, until the “right” place for a is found.

Implement the above algorithm for in-place, non-recursive insertion sort (remember that you do not need to `return` the sorted list; the modifications will be done directly to the input `l` and such modifications will be visible after the execution the function).

Solution. Here is a possible implementation:

```
def insertion_sort(l):
    for i in range(len(l)):
        j = i
        while j > 0 and l[j-1] > l[j]:
            l[j-1], l[j] = l[j], l[j-1]
            j -= 1
```

Looking at these 5 lines of code, you immediately see a further advantage of insertion sort: it is one of the simplest all-purpose sorting algorithms! (Compare it with the other algorithms). So, in spite of its poor efficiency, insertion sort is sometimes an interesting choice, especially when the lists to be sorted are not that big.

Exercise 4: quicksort with Lomuto's partitioning scheme

Lomuto's partitioning scheme is an in-place algorithm which takes a list l of length at least 2 and two positions $b < e$ within l , and returns a position $b \leq p \leq e$ while modifying l so that, at the end of the execution of the algorithm, it is partitioned as follows:

- for all $b \leq i < p$, $l[i] \leq l[p]$;
- for all $p < j \leq e$, $l[p] \leq l[j]$.

Here is the algorithm:

1. choose a pivot, place it in $l[e]$ and set $p = b$;
2. scan l from b to $e - 1$; if an element strictly smaller than the pivot is found, swap it with $l[p]$ and increment p ;
3. swap $l[p]$ and $l[e]$ and return p .

Implement it and then use it to implement quicksort in place. For simplicity, take the pivot to be $l[e]$ (so we avoid the initial swap at point 1).

The recursive function implementing quicksort will be of the form `qsortRec(l, b, e)`, where l is the list to sort and b and e delimit the portion where we are currently working. If p is the value returned by the partitioning function, it will contain a recursive call to `qsortRec(l, b, p-1)` and to `qsortRec(l, p+1, e)`. The actual quicksort function will then simply be

```
def quicksort(l):
    qsortRec(l, 0, len(l)-1)
```

Solution. Here is a possible implementation:

```
def partition(l, b, e):
    pivot = l[e]
    p = b
    for i in range(b, e):
        if l[i] < pivot:
            l[i], l[p] = l[p], l[i]
            p += 1
```

```

    l[e], l[p] = l[p], l[e]
    return p

def qsortRec(l, b, e):
    if b < e:
        p = partition(l, b, e)
        qsortRec(l, b, p-1)
        qsortRec(l, p+1, e)

def quicksort(l):
    qsortRec(l, 0, len(l)-1)

```

Exercise 5: quicksort with Hoare's partitioning scheme

Hoare's partitioning scheme is an in-place algorithm which takes a list l of length at least 2 and two positions $b < e$ within l , and returns a position $b \leq p < e$ (notice the difference with respect to Lomuto's scheme) while modifying l so that, at the end of the execution of the algorithm, it is partitioned as follows (again notice the difference with respect to Lomuto's scheme):

- for all $b \leq i \leq p$ and $p < j \leq e$, $l[i] \leq l[j]$.

Here is the algorithm:

1. choose a pivot;
2. set $i = b$, $j = e$;
3. increment i until $l[i] \geq \text{pivot}$;
4. decrement j until $l[j] \leq \text{pivot}$;
5. if $i \geq j$, go to 7;
6. otherwise, swap $l[i]$ and $l[j]$, increment i , decrement j and go to 3;
7. set $p = j$ or $p = e - 1$ in case $j = e$, and return p .

Implement it and then use it to implement quicksort in place. You may take the pivot to be the value of the middle element, that is, $\text{pivot} = l[b + (e-b+1)//2]$.

The recursive function implementing quicksort will be of the form $\text{qsortRec}(l, b, e)$, where l is the list to sort and b and e delimit the portion where we are currently working. If p is the value returned by the partitioning function, it will contain a recursive call to $\text{qsortRec}(l, b, p)$ and to $\text{qsortRec}(l, p+1, e)$ (notice the difference with respect to Lomuto's scheme). The actual quicksort function will be just like for Lomuto's scheme:

```

def quicksort(l):
    qsortRec(l, 0, len(l)-1)

```

Solution. Here is a possible implementation:

```

def partition(l, b, e):
    pivot = l[b+(e-b+1)//2]
    i = b-1
    j = e+1
    while i < j:
        if i >= b:
            l[i], l[j] = l[j], l[i]
        i += 1
        j -= 1
    while l[i] < pivot:
        i += 1
    while l[j] > pivot:
        j -= 1

```

```

    if j == e:
        j -= 1
    return j

def qsortRec(l, b, e):
    if b < e:
        p = partition(l, b, e)
        qsortRec(l, b, p)
        qsortRec(l, p+1, e)

def quicksort(l):
    qsortRec(l, 0, len(l)-1)

```

Exercise 6: counting sort

All sorting algorithms seen so far are *comparison-based*: they sort by making comparisons between elements of the list they are sorting. There are sorting algorithms that work without making any comparison. One such algorithm is *counting sort*. It works on input lists l whose values are assumed to range between 0 and $k - 1$, where k is some (usually small) positive constant. Here is a description:

- compute a list c of length k such that, for all $0 \leq i < k$, $c[i]$ is the number of occurrences of i in l ;
- the sorted list is now the list starting with $c[0]$ occurrences of 0, then $c[1]$ occurrences of 1, then $c[2]$ occurrences of 2 and so on until $c[k - 1]$ occurrences of $k - 1$.

1. Implement the above algorithm in place, that is, write the resulting list directly in l . For this implementation, fix $k = 100$.
2. Albeit it only works for lists containing integers whose size is bounded a priori (by k), counting sort has the advantage of being extremely fast. It may be useful even in the case of lists containing complex objects, of a priori unbounded size, as long as such objects have some sorting key of small size. For example, imagine having a list of records of the form

(FirstName, LastName, Weight, Height).

We may imagine this to be a list l whose elements are themselves lists of length 4, so that, for example, $l[i][2]$ is the weight of the i -th record in the list. Now, we may like to sort l according to *any* of the four fields: by last name, by first name, by weight or by height. Since the value of the latter two fields is bounded a priori,² it may be a good idea to apply counting sort for those.

In general, suppose that we are given a list of lists l , such that the j -th position in each list within l contains the sorting key, which is an integer between 0 and $k - 1$. We want to sort l ascendingly with respect to the j -th field, that is, if r is the sorted version of l , then for all positions $p \leq q$, we must have $r[p][j] \leq r[q][j]$, whereas there is no constraint on the fields other than j .

If we want to apply counting sort to this more complex situation, we cannot simply build the sorted list r out of the counting list c . Instead, for each element $l[i]$ of the input list, we need to compute the index in r where such an element will be located. This may be done by observing that, if $l[i][j] = h$, then certainly $l[i]$ will appear in r at a position greater or equal to $c[0] + c[1] + \dots + c[h - 1]$, because all positions before that will be occupied by elements of l whose j -th field contains a value strictly smaller than h . This suggests the following algorithm, which uses two further lists of length k , besides c .³

²Out of curiosity: the Guinness World Records apparently has it that the heaviest human being ever known weighed 442 kg, whereas the tallest human being ever known measured 272 cm. . .

³Actually, by cleverly reusing c , one may get away with only one list. But, for the purpose of understanding the algorithm, keeping three lists is conceptually simpler.

- a list b (the list of *base values*) such that $b[h]$ is the smallest position at which an element with key h may appear in r , computed as the above sum;
- a list o (the list of *offsets*) such that $o[h]$ acts like a counter, so that whenever we meet an element of l whose key is h , we will store it in $r[b[h] + o[h]]$, and increment $o[h]$.

Here is a description of the algorithm, which takes as input a list l containing n lists, a position j (the key field) and an integer k (the bound to the key values):

- compute a list c of length k such that, for all $0 \leq h < k$, $c[h]$ is the number of elements $l[i]$ of l such that $l[i][j] = h$;
- compute a list b of length k such that, for all $0 \leq h < k$, $b[h] = \sum_{i=0}^{h-1} c[i]$;
- initialize a list o containing k zeros;
- initialize the output list r , of length n ;
- for each $0 \leq i < n$, let $h = l[i][j]$, copy $l[i]$ in $r[b[h] + o[h]]$ and increment $o[h]$;
- return r .

Implement the above algorithm as a function `csort(l, j, k)`. Notice that this will *not* be in place. NB: for building the list b , you are *not* allowed to use Python's list slicing and the built-in `sum` function. In fact, there's a simple way of doing this which does not need a sum function at all.

Solution. Here are possible implementations:

1.

```
def countingsort(l):
    c = [0] * 100
    n = len(l)
    for i in range(n):
        c[l[i]] += 1
    i = 0
    for k in range(100):
        for _ in range(c[k]):
            l[i] = k
            i += 1
```

2.

```
def csort(l, j, k):
    c = [0] * k
    n = len(l)
    for i in range(n):
        c[l[i][j]] += 1
    b = [0] * k
    for i in range(1, k):
        b[i] = b[i-1] + c[i-1]
    o = [0] * k
    r = [0] * n
    for i in range(n):
        v = l[i][j]
        r[b[v] + o[v]] = l[i]
        o[v] += 1
    return r
```

Exercise 7: radix sort

As mentioned above, counting sort is very fast, but it only works for sorting lists of small integers. Nevertheless, it may be used as the building block of another non-comparison-based algorithm, called *radix sort*. Radix sort is fully general, that is, it may be applied to lists containing integers whose value is not bounded a priori. It is also very efficient: its worst-case asymptotic complexity is the same as that of merge sort, but usually it is faster and in fact its performance may rival that of quicksort. Moreover, when the size of the integers contained in the input list l is much smaller than the length of l , radix sort performs better than any comparison-based algorithm.

Let us first consider radix sort in its general form, which is for sorting lists according to the *lexicographic order*. Two lists l_1, l_2 of comparable items and of equal length are *lexicographically ordered* as $l_1 \leq l_2$ if they are either equal, or they first differ at a position h such that $l_1[h] < l_2[h]$. The alphabetical order is an example of lexicographic order: bounded < bourbon.

Radix sort uses counting sort to solve the following problem:

LEXICOGRAPHIC SORT:

input: a list l of lists of fixed length m , containing integers between 0 and $k - 1$;

output: the list l lexicographically ordered, ascendingly.

The algorithm is extremely simple: repeatedly apply counting sort to l , starting with $m - 1$ as the key (the “least significant field”), then $m - 2$ and so on until 0 (the “most significant field”). For example, if the input list is

$[[2, 1, 2], [1, 1, 0], [2, 2, 0], [0, 1, 1], [1, 1, 1]],$

then sorting with 2 (the last field) as the key yields

$[[1, 1, 0], [2, 2, 0], [0, 1, 1], [1, 1, 1], [2, 1, 2]],$

from which, sorting with 1 (the middle field) as the key, we get

$[[1, 1, 0], [0, 1, 1], [1, 1, 1], [2, 1, 2], [2, 2, 0]]$

and then, sorting with 0 (the first field) as the key, we finally obtain

$[[0, 1, 1], [1, 1, 0], [1, 1, 1], [2, 1, 2], [2, 2, 0]],$

which is the input list sorted lexicographically. Notice that, if we interpret the elements of the inner lists as digits, and the lists as numbers (for example, $[1, 1, 0]$ as the number 110), and if we ignore leading zeroes, then the final list corresponds to the numerically sorted list $[11, 110, 111, 212, 220]$. This is because, after adding leading zeroes, the usual numerical order coincides with the lexicographic order on digit expansions (with respect to any base).

Implement the above algorithm via a function `lexsort(l, m, k)`, using the function `csort(l, j, k)` you wrote at the end of Exercise 6. Then, write the following functions:

- `maximum(l)`, which takes a list of integers and returns its maximum value. NB: this already exists in Python as the built-in function `max`, which you are of course *not* allowed to use; you are asked to implement it yourself.
- `declen(n)`, which takes a non-negative integer n and returns the number of digits in the decimal expansion of n (for example, `declen(107)` returns 3). NB: you are *not* allowed to use Python’s built-in logarithm function for this!
- `decexp(n, m)`, which takes two non-negative integers n and m and returns a list l of length m such that $l[i]$ is the $(m - i)$ -th decimal digit of n , from the right, with leading zeros if the decimal expansion of n has less than m digits (for example, `decexp(1203, 3)` returns $[2, 0, 3]$, whereas `decexp(54, 3)` returns $[0, 5, 4]$). NB: you are *not* allowed to use Python’s integer-to-string conversion and then play with list slicing for this!
- `deccmp(l)`, which takes a list of integers between 0 and 9 and returns the number such that l is its decimal expansion (for example, `deccmp([0, 1, 2])` returns 12).

(To test your implementations, you may check that, for any list of digits l , `decexp(deccmp(l), len(l))` returns l itself, and for any non-negative integer n , `deccmp(decexp(n, declen(n)))` returns n itself). Using the above functions, implement radix sort on non-negative integer lists as follows:

- find the maximum value mx of the input list l , and let $m = \text{declen}(mx)$;
- using `decexp`, convert l into a list of lists of length m ;
- apply `lexsort` to sort this list lexicographically, obtaining a list of lists r ;
- convert r into a list of non-negative integers by applying `deccmp`;
- return r .

Solution. Here are possible implementations (we recall the `csort` function from Exercise 6, just so that the code below will run stand-alone):

This is the `csort` function from the previous exercise

```
def csort(l,j,k):
    c = [0] * k
    n = len(l)
    for i in range(n):
        c[l[i][j]] += 1
    b = [0] * k
    for i in range(1,k):
        b[i] = b[i-1] + c[i-1]
    o = [0] * k
    r = [0] * n
    for i in range(n):
        v = l[i][j]
        r[b[v] + o[v]] = l[i]
        o[v] += 1
    return r
```

```
def lexsort(l,m,k):
    for j in range(m):
        l = csort(l,m-j-1,k)
    return l
```

```
def declen(n):
    if n == 0:
        return 1
    l = 0
    while n > 0:
        n = n // 10
        l += 1
    return l
```

```
def decexp(n,k):
    l = [0] * k
    for i in range(k):
        l[k-i-1] = n % 10
        n = n // 10
    return l
```

```
def deccmp(l):
    b = 1
    n = 0
    m = len(l)
    for i in range(m):
        n += l[m-i-1] * b
```

```

        b *= 10
    return n

def radixsort(l):
    if l == []:
        return l
    mx = l[0]
    n = len(l)
    for i in range(n):
        if l[i] > mx:
            mx = l[i]
    m = declen(mx)
    d = [0] * n
    for i in range(n):
        d[i] = decexp(l[i], m)
    d = lexsort(d, m, 10)
    for i in range(n):
        l[i] = deccomp(d[i])

```

Exercise 8: stability

Consider again a list of records of the form

(FirstName, LastName, Weight, Height).

As we said in Exercise 6, we may be interested in sorting this list according to any field. Our current implementations of comparison-based algorithms (merge sort, insertion sort, quicksort) do not allow this, because they consider the records as a whole (in fact, in this situation they would sort according to lexicographic order, see point 4 of Exercise 9).

Modify each of these algorithms (merge sort, insertion sort, quicksort with both partition schemes) so that they take as input a list of lists l and an integer k , the *sorting key*, and they sort according to the value of the k -th field of the inner lists. For example, if

$$l = [[1, 2, 5], [0, 3, 6], [5, 8, 2], [2, 1, 9], [8, 8, 7]],$$

then on input l and with $k = 1$ the sorted list could be

$$[[2, 1, 9], [1, 2, 5], [0, 3, 6], [5, 8, 2], [8, 8, 7]]$$

(we are using the second field as key, and we are sorting ascendingly with respect to it). (NB: *this modification should be trivial! It should simply take copy-and-pasting your code, adding an argument k to your functions and appending $[k]$ in a few places here and there. If you are taking too long to do this, something's not right...*).

A sorting algorithm is said to be *stable* if it preserves the original order of elements with the same sorting key. In the above example, where the sorting key is $k = 1$, the lists $[5, 8, 2]$ and $[8, 8, 7]$ have the same sorting key (the second field is 8 in both), so the list

$$[[2, 1, 9], [1, 2, 5], [0, 3, 6], [8, 8, 7], [5, 8, 2]]$$

is also sorted according to the second field and is a perfectly acceptable result (that is why we said “the sorted list could be” and not “should be”). However, a sorting algorithm returning the above list is *not* stable, because, in the original list, $[8, 8, 7]$ came *after* $[5, 8, 2]$, whereas now their relative order is reversed.

Stability is interesting when one wants to order relatively to a main key and a secondary key. For instance, imagine that we want to sort the above records alphabetically by last name, and then, if

more people have the same last name, we want them to appear in alphabetical order with respect to their first names. We may achieve this with two sorting passes, as long as we use a stable sorting algorithm for the second pass: the first one sorts alphabetically by first name, the second by last name. If the algorithm used for the second pass is not stable, there is no guarantee that the first-name-based ordering between records with the same last name will be respected.

Take the modified version “with key” (that you just wrote for this exercise) of merge sort, insertion sort, quicksort (both variants), as well as the function `csort` implementing counting sort “with key” that you wrote at the end of Exercise 6, and test it on a list of the form

```
l = [[random.randint(0,9),i] for i in range(20)]
```

using as sorting key $k = 0$ (the first field, which has small random values). Since the second field coincides with the position of the element in the original list, you will be able to quickly assess the stability of the result. Which algorithm seems to be stable and which one does not? Would radix sort work if `csort` were not stable?

Solution. Here are versions “with key” of all the required sorting algorithms:

```
# Merge sort with key
def mergekey(l1,l2,k):
    n1 = len(l1)
    n2 = len(l2)
    l = [0] * (n1 + n2)
    i1 = 0
    i2 = 0
    i = 0
    while i1 < n1 and i2 < n2:
        if l1[i1][k] <= l2[i2][k]:
            l[i] = l1[i1]
            i1 += 1
        else:
            l[i] = l2[i2]
            i2 += 1
        i += 1
    if i1 < n1:
        for i1 in range(i1,n1):
            l[i] = l1[i1]
            i += 1
    elif i2 < n2:
        for i2 in range(i2,n2):
            l[i] = l2[i2]
            i += 1
    return l

def mergesortkey(l,k):
    n = len(l)
    if n <= 1:
        return l
    h = n // 2
    l1 = [0] * h
    l2 = [0] * (n-h)
    for i in range(n):
        if i < h:
            l1[i] = l[i]
```

```

        else:
            l2[i-h] = l[i]
        return mergekey(mergesortkey(l1,k),mergesortkey(l2,k),k)

# Insertion sort with key
def insertionsortkey(l,k):
    for i in range(len(l)):
        j = i
        while j > 0 and l[j-1][k] > l[j][k]:
            l[j-1],l[j] = l[j],l[j-1]
            j -= 1

# Lomuto quicksort with key
def LomutoKey(l,b,e,k):
    pivot = l[e][k]
    p = b
    for i in range(b,e):
        if l[i][k] < pivot:
            l[i],l[p] = l[p],l[i]
            p += 1
    l[e],l[p] = l[p],l[e]
    return p

def qsortlRec(l,b,e,k):
    if b < e:
        p = LomutoKey(l,b,e,k)
        qsortlRec(l,b,p-1,k)
        qsortlRec(l,p+1,e,k)

def qsortLomutoKey(l,k):
    qsortlRec(l,0,len(l)-1,k)

# Hoare quicksort with key
def HoareKey(l,b,e,k):
    pivot = l[b+(e-b+1)//2][k]
    i = b-1
    j = e+1
    while i < j:
        if i >= b:
            l[i],l[j] = l[j],l[i]
            i += 1
            j -= 1
        while l[i][k] < pivot:
            i += 1
        while l[j][k] > pivot:
            j -= 1
    if j == e:
        j -= 1
    return j

```

```
def qsorthRec(l, b, e, k):
    if b < e:
        p = HoareKey(l, b, e, k)
        qsorthRec(l, b, p, k)
        qsorthRec(l, p+1, e, k)

def qsortHoareKey(l, k):
    qsorthRec(l, 0, len(l)-1, k)
```

(Notice how it's just a copy-and-paste of the solution to the relevant exercise, with some [k] suitably added).

For what concerns stability, the only non-stable sorting algorithm here is quicksort, in both of its variants.

The stability of counting sort of course is what allows its use in implementing radix sort: lexicographic sorting may be achieved by repeatedly sorting with respect to each key, from the “least significant” to the “most significant”, only if each pass preserves the ordering established by the previous pass, which is possible only if the sorting algorithm used is stable.

Exercise 9: experiment!

Now that you have programmed lots of different sorting algorithms, you may test their speed! For instance, try the following:

1. use `import random` to enable invoking Python's random number generator, then generate a random list of 5000 integers in the range 0–999 with the instruction

```
l = [random.randint(0,999) for _ in range(5000)]
```

Apply all the different sorting algorithms to random lists like this (except counting sort, which would not work because the elements go beyond 99). You should see that all algorithms work instantaneously, except one... which one?

2. In order to test the difference between the “fast” algorithms, you'll need to time their execution. Use `import time` to enable access to the internal clock. Then, you may time the execution of a program as follows:

```
start = time.time()
# execute something
t = time.time() - start
print("The execution took %.6f seconds" %t)
```

The last instruction will print the time up to microseconds. Now try and time the “fast” algorithms (merge sort, the two versions of quicksort, radix sort), maybe using even longer lists, but with smaller integers so we may also test counting sort:

```
l = [random.randint(0,99) for _ in range(10000)]
```

or even

```
l = [random.randint(0,9) for _ in range(10000)]
```

Who's the winner of the “contest”? Do you see any difference between the two versions of quicksort?

3. Now, instead of testing the algorithms on random lists, let's test them on “special” lists. For example, a constant list of length 2000:

```
l = [5] * 2000
```

or an already sorted list of the same length:

```
l = [i for i in range(2000)]
```

How does quicksort with Lomuto's partitioning scheme behave on such lists? Is its performance closer to merge sort or to insertion sort?

- As mentioned in Exercise 7, the original application of radix sort is for sorting lexicographically. It happens that Python already has lexicographical ordering built in: if `l1` and `l2` are two integer lists, writing for example `l1 < l2` will test whether `l1` is strictly smaller than `l2` in the lexicographic order. Any other comparison operator may also be applied (`>`, `>=`, `<=`). This means that your Python implementation of comparison-based algorithms (merge sort, quicksort, insertion sort) automatically works for lists of integer lists, and sorts them ascendingly according to the lexicographic order. Therefore, we may compare the performance of these functions with the function `lexsort` you wrote in Exercise 7. Try executing merge sort, quicksort (with Hoare's scheme, Lomuto's will do poorly here) and `lexsort(l, 3, 10)` when `l` is a random list of triples

```
l = [[random.randint(0,9), random.randint(0,9), random.randint(0,9)]
      for _ in range(n)]
```

with `n` progressively growing from 10^3 to 10^6 . What do you notice?

Solution.

- Insertion sort is the slow algorithm. Indeed, its complexity is quadratic (like selection sort), as opposed to the other algorithms we consider, which all have average-case complexity $n \log n$.
- Counting sort should be by far the fastest. Quicksort with Hoare's partitioning scheme should come in second, with radix sort following closely behind. On random lists with large integers, Lomuto's partitioning scheme rivals with Hoare's in efficiency (in fact, in many cases it is even faster), but as soon as the integer size becomes small with respect to the length of the input list, its performance degenerates and becomes much worse than that of merge sort.
- Lomuto's partitioning scheme behaves very badly with already-sorted lists and with constant lists: it degenerates to insertion sort (in fact, since its implementation is recursive and more complex, it is even worse in practice!). The other algorithm's performance should not differ too much with respect to the case of random lists.
- Radix sort (that is, the function `lexsort`) should be by far the fastest: its complexity is linear in $n!$. Quicksort and merge sort both have complexity $n \log n$, but merge sort has bigger constants, so its performance is not as good. The $\log n$ factor should show up very clearly in the comparison between the running time of quicksort and radix sort.